


Part 1: RDBMS and NOSQL

Valerie Hayot-Sasson


Why distribute data ?

Size




Data too large for a single node

Performance



Satisfy the needs of many clients

Reliability




Redundancy in case of failure

Desirable properties for distributing data

Consistency : All entities see the **same data**




Availability: Every request received a **non-error response for every operation**



Partition tolerance: The system **continues to operate** even when connections between nodes are down/slow

The CAP Theorem

A network-shared data system can only have **two** out of the three desirable properties



Eric Brewer

Takeaways

- Network partitions are inevitable: must choose between C and A
 - Forfeit consistency: process the operation even though it might be wrong/lead to inconsistency
 - Forfeit availability: return an error if we cannot guarantee that data is up to data
- Properties are continuous: do not need to forego one all the time
- Network partitions are rare: only make tradeoff when partitions occur

[\[Brewer, E. \(2012\). CAP twelve years later: How the "rules" have changed. Computer, 45\(2\), 23-29.\]](#)

Consistency versus Availability

Databases (C over A)

- **Atomic**
- **Consistent***
- **Isolation**
- **Durability**

NoSQL (A over C)

- **Basically Available**
- **Soft state**
- **Eventual consistency**

* Consistency in databases includes preserving all databases rules. CAP's consistency involves single-copy consistency

RDBMS to NoSQL

Many current workloads do not require ACID properties, allowing us to loosen consistency requirements (eventual consistency) to improve performance and availability

- Benefits: simple scale out, faster read/writes, maybe easier to use?

Some slides based on: <https://courses.engr.illinois.edu/cs425/fa2016/L9-11.FA16.pdf>

Relational Database Management Systems (RDBMS)

- Schema oriented
- Data organized into tables
 - Defined columns
 - Constraints (e.g., foreign key)
 - Types
- Supports joins between tables
- Structured Query Language (SQL)
- Highly Optimized Query Planners

ID	Name	Email
1	Jane	jane@
2	Bob	bob@
3	Mary	mary@

```
SELECT *  
FROM USERS  
WHERE name = 'Mary'
```

ID	Name	Email
3	Mary	mary@

grade_id	homework	student_id	grade
1	1	1	60
2	1	2	40
3	2	1	75

```
SELECT users.name, grades.grade  
FROM users JOIN grades  
ON users.id = grades.student_id
```

name	grade
Jane	60
Bob	40
Jane	75

ACID Properties


- **Atomic:** Transactions execute in “all-or-nothing” manner
 - Transactions can be one or more operations
 - Example: buying a plane ticket
- **Consistency:** Transition from one valid state to another.
 - All DB rules must be preserved following a transaction.
- **Isolation:** No transaction sees uncommitted changes of concurrent transactions.
 - Read phenomena: dirty reads, non-repeatable reads, phantom reads
- **Durability:** State changes caused by committed transactions are preserved.

Scaling RDBMS

- ACID properties make scaling out difficult. More common to scale up (add more resources).
- Distribution approach depends on needs:
 - **Read heavy:** add replicas
 - Add as many replicas needed to handle the workload
 - Balance queries across replicas in Round-Robin fashion
 - **Write heavy:** Partition database across many servers
 - Put different tables on different machines or shard one table across many machines
 - Horizontal fragmentation: split selection query
 - Vertical fragmentation: split by projection queries

Horizontal Fragmentation

id	Name	Sex	Age
1	Sebastian	M	3
2	John	M	5
3	Garry	M	8
4	Louis	F	7



Vertical Fragmentation

id	Name	Sex	Age
1	Sebastian	M	3
2	John	M	5
3	Garry	M	8
4	Louis	F	7


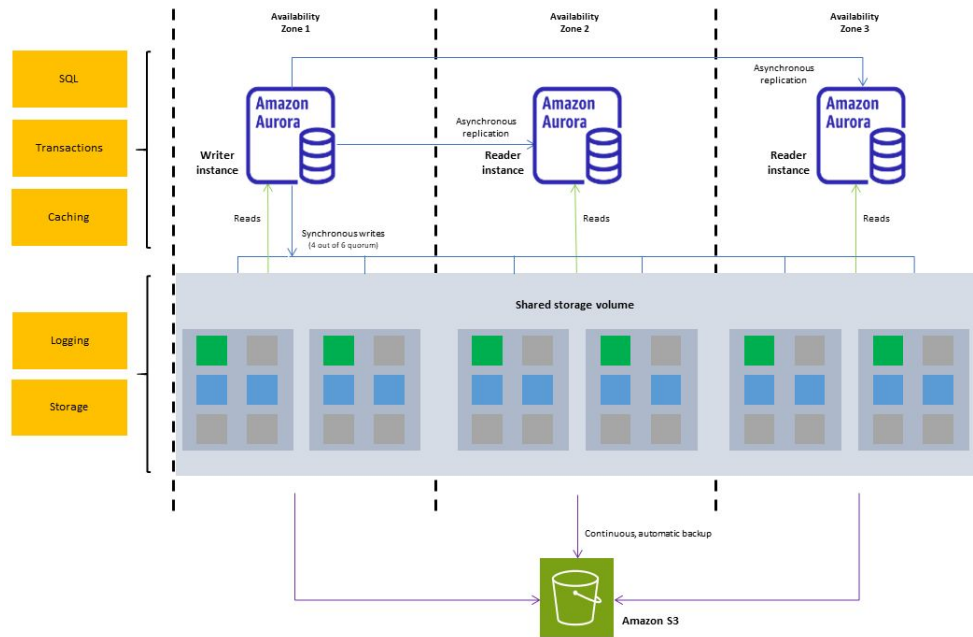


Image credit:

<https://stackoverflow.com/a/61680342>

Amazon Aurora

MySQL and PostgreSQL-compatible relational database in the cloud



Adapting to evolving workload needs

- Newer workloads:
 - Data are large and unstructured
 - Random access read/writes
 - Joins are infrequent
- NoSQL Goals:
 - Speed
 - Avoid single point of failure
 - Reduce ownership costs
 - Fewer system administrators
 - Elastic scalability
 - Scale out

Not Only SQL (NoSQL)

- Unstructured
- No schema
- Sparse columns/rows
- No foreign keys (joins might not be supported)
- Does not use SQL as a query language

1	Name: kyle, grade: 100
2	Name: bob, grade: 60
3	Name: jane, grade: 95

ID	Document
1	Name: Kyle
2	Name: bob, email: bob@
3	Email: jane@


NoSQL advertised benefits

- Simplicity (from a design/interaction perspective)
- Horizontal scaling
- High availability
- Performance (depending on the use case)

Rows vs Columns


- Row store: Store all attributes of a tuple together

- Advantage: easy to add/modify records
- Disadvantage: may read unnecessary data



→ Column store: Store all rows of for an attribute together

- ◆ Advantage: only read relevant attributes
- ◆ Read/write requires multiple accesses



NoSQL variations

- Key-value stores
 - Store simple, flat lists of key-value pairs
 - Expose get and put operations
 - Redis, Memcached
- Document stores
 - Each key is paired with a whole document (e.g., JSON)
 - MongoDB, CouchDB
- Column family
 - Each row is addressed by a key and contains one or more columns
 - Columns are key-value pairs (tuple: name, value, timestamp)
 - Cassandra, HBase
- Graph
 - Neo4J

Organizing NoSQL stores


- Most NoSQL stores group key-value pairs into something that looks like a table
 - Cassandra: Column families
 - MongoDB: Collections
 - HBase: Tables
 - DynamoDB: Tables
- NoSQL stores typically store a column (or group of columns) together
 - Index entries in a column for easy retrieval
 - Supports range queries so you do not need to get entire database

Eventual consistency

- Tradeoff between consistency and availability when there is a network partition
 - RDBMS: strong consistency over availability
 - NoSQL: availability over consistency (eventual consistency)
- If writes to a key stop, then all values (replicas) will eventually converge
 - If writes continue, we will keep trying to converge
 - Maintain set of updated values lagging latest values sent by clients
- Implication: may return stale values to clients (e.g., during many back-to-back writes)
- Work well with periods of low writes – system converges quickly

Tunable consistency levels

- Client is allowed to choose consistency level for each operation (read/write)
- Strong consistency: $R + W > N$
- Eventual consistency: $R + W \leq N$
 - ANY: write written to at least one node
 - **Fastest**: node will cache write and reply to client quickly
 - ALL: write written to all nodes
 - Ensures strong consistency, but **slow**
 - ONE: At least one replica
 - Faster than ALL but cannot tolerate more than one failure
 - QUORUM: written to a configurable quorum across nodes



<https://docs.datastax.com/en/cassandra/3.0/cassandra/dml/dmlConfigConsistency.html>



- Created in 2007
- Document-oriented database
 - Schema-less
 - Stores arbitrary documents
 - Documents are given a unique ID (`_id`)
 - Documents contain arbitrary keys/values
 - Provides its own query language
 - Leader-worker replication across replica sets
 - No joins/transactions

MongoDB architecture

- Data is sharded across machines
 - Scale and performance
- Documents are distributed across shards using an immutable “shard key” (from keys in documents)
 - Knowledge of data/queries can optimize use
- Supports different sharding strategies
 - E.g., Hashing and range-based
- Query router sends queries to appropriate shards

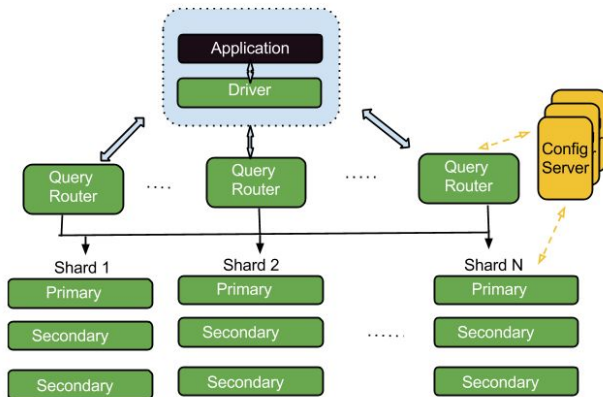


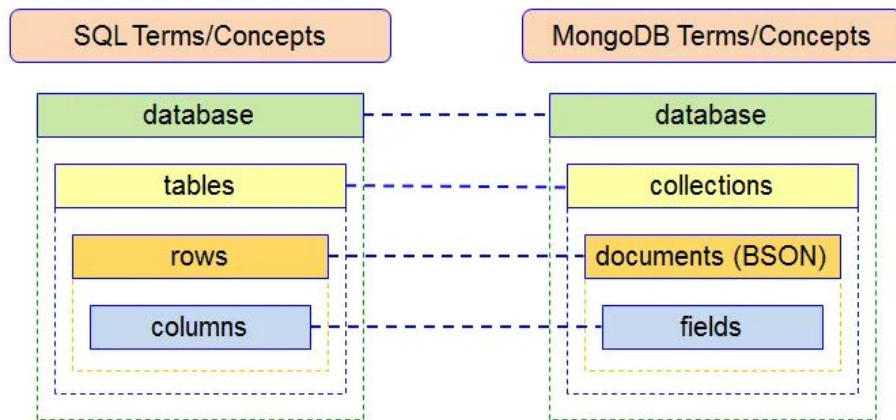
Image credit:

<https://medium.datadriveninvestor.com/mastering-mongodb-understanding-and-utilizing-the-nosql-database-41ec3890f97>

BSON

- Mongo wanted to support users and JSON is pervasive
- Challenges of JSON:
 - Limited data types (no binary, date)
 - Objects do not have fixed length (slows processing)
- BSON (Binary JSON)
 - Adds data types and fixed lengths
 - Anything in JSON can be stored as BSON
- MongoDB stores data in BSON internally and over the network

MongoDB concepts



<https://medium.com/nerd-for-tech/all-basics-of-mongodb-in-10-minutes-baddaf6b6625>

MongoDB exercise

- Installation

- a. Docker database server

- ```
docker run -p 27017:27017 mongo:latest
```

- b. Python client (in your terminal)

- ```
pip install pymongo
```

- Docs: <https://pymongo.readthedocs.io/en/stable/>

Create database/column

```
import pymongo
mongo_client = pymongo.MongoClient("mongodb://localhost:27017/")
db = mongo_client["mydatabase"]
col = db["customers"]

# Insert records
l = [{"name": "Kyle", "office": "303"},
     {"name": "Tyler", "office": "312"}, ]

x = col.insert_many(l)
```

Query things

Search for “kyle”

```
for x in col.find({'name': 'Kyle'}):  
    print(x)
```

Get only some attributes

```
for x in col.find({}, { "_id": 0, "office": 1 }):  
    print(x)
```

Exercise

1. Create a database that stores students and their grades for MPCs course
 - E.g., `student_name`, `student_id`, `hw1`, `hw2`, `hw3`
2. Insert 5+ records
3. Find number of total records (hint: `collection count_documents`)
4. Find a student by name and get only their name and hw1 grades
5. Find students who received more than 50 on hw1 (read the docs for query syntax)
6. (Optional/advanced) Compute each student's grade for the course

Submission

https://classroom.github.com/a/Z8_TBszc